

# Neural Networks

**Jingwei Tang @ CGL**

01.04.2019 – 05.04.2019

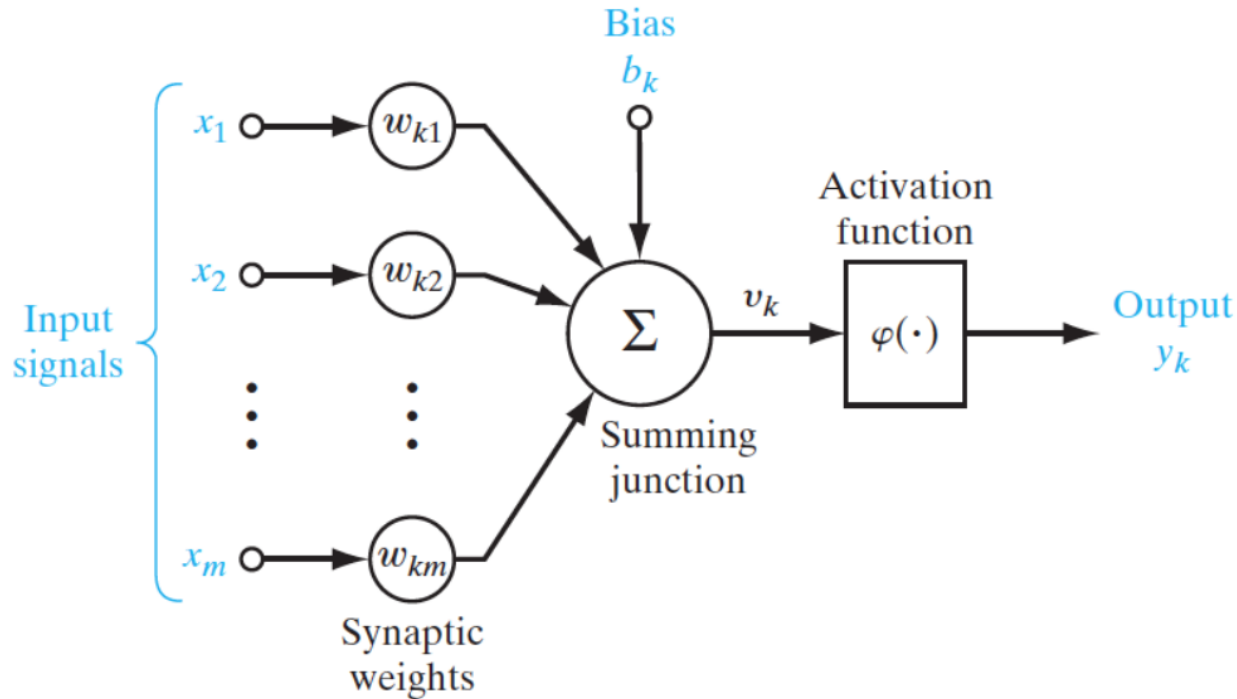
[jingwei.tang@inf.ethz.ch](mailto:jingwei.tang@inf.ethz.ch)

**1. Recap of ANN architectures.**

2. Solving One Exam Question.

3. PyTorch Demos.

# Neurons



Haykin, Simon S., et al. Neural networks and learning machines. Vol. 3. Upper Saddle River: Pearson, 2009.

**Linear Units:**

$$g: \mathbb{R}_m^m \rightarrow \mathbb{R},$$
$$g(\mathbf{x}) = \sum_{j=1}^m w_j x_j + b$$

**Activation function:**

$$\phi: \mathbb{R} \rightarrow \mathbb{R}$$

**Output:**


$$y = \phi(g(\mathbf{x}))$$

# Activation Functions

- **Sigmoid Units:**

Logistic function:  $\sigma(z) = \frac{1}{1+e^{-z}} \in (0,1)$

Tanh function:  $\tanh(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}} \in (-1,1)$



# Activation Functions


- **Rectified Linear Unit (ReLU):**

$$(z)_+ = \max(0, z)$$

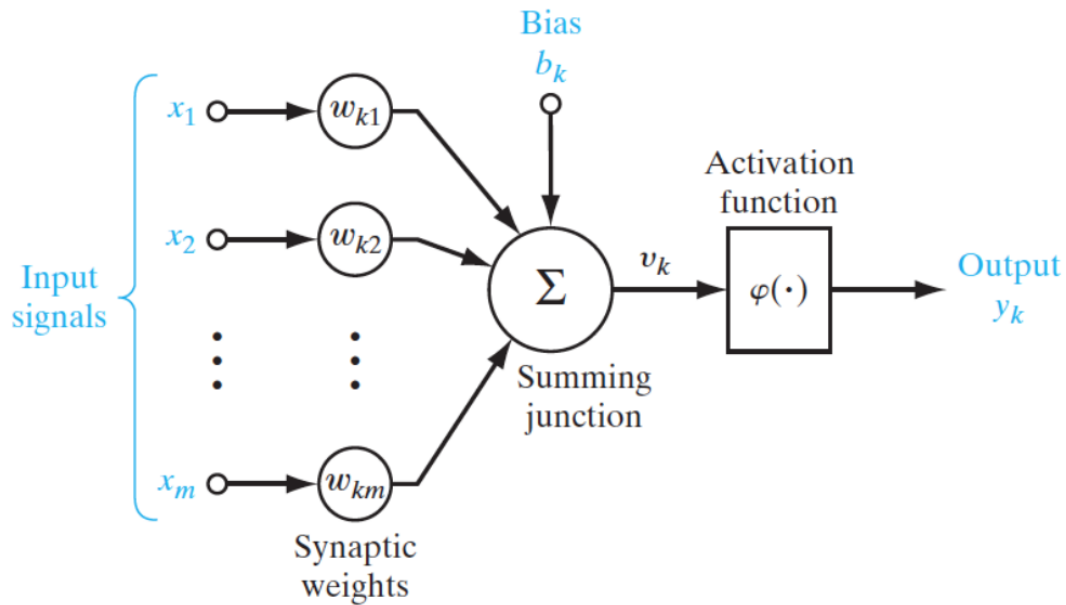
- **Leaky ReLU:**

$$(z)_{+, \alpha} = \begin{cases} \alpha z, & z < 0 \\ z, & z \geq 0 \end{cases}$$

With  $\alpha$  being a small value (0.001).



# Single Neuron as Binary Classifier



- **Binary Softmax classifier:**

- Use sigmoid activation function.
- Interpret  $\sigma(\sum_j w_j x_j + b)$  to be  $P(y = 1|x; w)$ .

- **Binary SVM classifier:**

- Use an extra max-margin hinge loss to the output.


# Layers of Neurons (Fully Connected Layers)

- **Single Hidden layer**

$$F(\mathbf{x}) = \sum_{i=1}^k w_i^{(2)} \phi \left( \sum_{j=1}^m w_{ij}^{(1)} x_j \right)$$

Compact representation:

$$F(\mathbf{x}) = \mathbf{W}^{(2)} \phi(\mathbf{W}^{(1)} \mathbf{x})$$




[https://en.wikipedia.org/wiki/Neural\\_network](https://en.wikipedia.org/wiki/Neural_network)

# Layers of Neurons (Fully Connected Layers)

- **Deep Multi-Layer Networks (ANN, MLP):**

$$F(\mathbf{x}) = \phi^{(L)}(W^{(L)} \phi^{(L-1)}(W^{(L-1)} \dots \phi^{(1)}(W^{(1)} \mathbf{x}) \dots))$$





# Multi-layer Network as Regressor

- **Output:**

Real-valued output neuron(s), without activation function.

- **Loss function:**

$\mathcal{L}_1$  or  $\mathcal{L}_2$  distances between predicted and ground-truth output.

$$l(\mathbf{y}^*, \mathbf{y}) = \|\mathbf{y}^* - \mathbf{y}\|_1$$

# Multi-layer Network as Classifier

## Binary:

- **Output:** Single output neuron  $y$ , use **sigmoid** activation function as probability of class membership.

$$\sigma = \frac{1}{1 + e^{-y}}$$

- **Loss:** Use **(binary) cross-entropy loss**. ( $y^* \in \{0, 1\}$ )  
$$l(y^*, y) = -y^* \log(\sigma) - (1 - y^*) \log(1 - \sigma)$$

# Multi-layer Network as Classifier

## Multi-class ( $C$ classes):

- **Output:** Multiple output neurons  $\mathbf{y}$ , use **SOFTMAX** function. Interpret output values as probability for corresponding class.

$$\sigma_i = \frac{e^{y_i}}{\sum_j e^{y_j}}$$

- **Loss:** Use **cross-entropy loss**. ( $\mathbf{y}^*$  is encoded as one-hot vector with one positive class (1) and  $C - 1$  negative classes (0).)

$$l(\mathbf{y}^*, \mathbf{y}) = - \sum_i y_i^* \log(\sigma_i)$$


# Universal Approximation Theorem

- Given **ANY** continuous function  $f(x)$  and some  $\epsilon > 0$ , there exists a Neural Network  $g(x)$  with one hidden layer (with a reasonable choice of **non-linearity**, e.g. sigmoid) such that

$$\forall x, |f(x) - g(x)| < \epsilon.$$

Complete proof in [Approximation by Superpositions of Sigmoidal Functions \(1989\)](#).


[Intuitive explanation](#) from Michael Nielson.



1. Recap of ANN architectures.
- 2. Solving One Exam Question.**
3. PyTorch Demos.

## Exam 2016 Question 5

Consider the following neural network with two logistic hidden units  $h_1, h_2$ , and three inputs  $x_1, x_2, x_3$ . The output neuron  $f$  is a linear unit, and we are using the squared error cost function  $E = (y - f)^2$ . The logistic function is defined as  $\rho(x) = 1 / (1 + e^{-x})$ .



- (i) Consider a single training example  $\mathbf{x} = [x_1, x_2, x_3]$  with target output (label)  $y$ . Write down the sequence of calculations required to compute the squared error cost (called forward propagation).
- (ii) A way to reduce the number of parameters to avoid overfitting is to tie certain weights together, so that they share a parameter. Suppose we decide to tie the weights  $w_1$  and  $w_4$ , so that  $w_1 = w_4 = w_{\text{tied}}$ . What is the derivative of the error  $E$  with respect to  $w_{\text{tied}}$ , i.e.  $\nabla_{w_{\text{tied}}} E$ ?

(i)

$$h_1 = \phi(w_1x_1 + w_3x_2 + w_5x_3)$$

$$h_2 = \phi(w_2x_1 + w_4x_2 + w_6x_3)$$

$$f = u_1h_1 + u_2h_2$$


$$E = (y - f)^2$$

(ii)

$$\frac{\partial E}{\partial w_{tied}} = \frac{\partial E}{\partial f} \left( \frac{\partial f}{\partial h_1} \frac{\partial h_1}{\partial w_{tied}} + \frac{\partial f}{\partial h_2} \frac{\partial h_2}{\partial w_{tied}} \right)$$

$$\begin{aligned} \frac{\partial E}{\partial w_{tied}} = & -2(y - f)(u_1x_1\phi'(w_1x_1 + w_3x_2 + w_5x_3) \\ & + u_2x_2\phi'(w_2x_1 + w_4x_2 + w_6x_3)) \end{aligned}$$

$$\frac{\partial E}{\partial w_{tied}} = -2(y - f)(u_1x_1h_1(1 - h_1) + u_2x_2h_2(1 - h_2))$$



(iii) For a data set  $D = \{(\mathbf{x}^{(1)}, y^{(1)}), \dots, (\mathbf{x}^{(n)}, y^{(n)})\}$  consisting of  $n$  labeled examples, augment the pseudocode of the stochastic gradient descent algorithm below with learning rate  $\eta_t$  for optimizing the weight  $w_{\text{tied}}$  (assume all the other parameters are fixed).

**begin**

$w_{\text{tied}} \leftarrow 0, \eta_t = 1/t;$

**for**  $t = 1$  *to*  $T$  **do**

// Fill in code to implement SGD

Select  $(x^{(i)}, y^{(i)})$  from  $D$  uniformly at random.

Forward pass like in (i) to compute  $h_1, h_2, f, E$ .

Backward pass like in (ii) to compute  $\frac{\partial E}{\partial w_{\text{tied}}}$ .

Update parameter  $w_{\text{tied}} = w_{\text{tied}} - \eta_t \frac{\partial E}{\partial w_{\text{tied}}}$ .

**end**

**end**



1. Recap of ANN architectures.
2. Solving One Exam Question.
3. **PyTorch Demos.**

# References

- [Stanford CS231n](#)
- [Coursera Deep Learning Specialization](#)